

What is a Support Vector Machine (SVM)?

What is a Support Vector Machine (SVM)?

- Imagine you have two types of fruits: *apples* and *oranges* mixed together on a table.
- You want to draw *a straight line* on the table so all *apples* are on one side and all *oranges* on the other.
- That is what a SVM does: *it draws the best straight line to separate two groups.*

Why “Best” Line?

- *Not Just Any Line*
- There can be many lines that separate apples and oranges.
- SVM chooses the one that **leaves the biggest space** (margin) between the line and the closest **apple** or **orange**.
- **Bigger space** = more confidence the line will still work for new fruits.


Why “Best” Line?

Not Just Any Line

There can be many lines that separate apples and oranges.

SVM chooses the one that leaves the biggest space (**margin**) between the line and the closest apple or orange

biggest space
(margin)




Stephen and Peace

- *Blue line is the best that can separate apples and oranges.*

How it works?

- SVM imagines a boundary that's far from both groups.
- The closest points (one from each group) are called “support vectors” , they are like referees marking the edge of the field.
- The SVM positions the line so it's **as far from those referees as possible**.

Why “Best” Line?



SVM algorithm

- **Measure two things** about each fruit — for example, weight and color shade.
- **Plot them on a chart:** apples cluster in one area, oranges in another.
- **SVM's job:** Find the best **straight line** (or curve) that divides apples from oranges.
- **Maximize the gap** between this line and the nearest fruit from each side — this gap is called the **margin**.
- **Support vectors** are the fruits that sit right at the edge of that margin — they “support” the line's position.
- When a **new fruit** comes in, SVM checks which side of the line it falls on to decide if it's an apple or an orange.

SVM classification for Orange vs Apple


- **SVM decision rule**

- In a **linear SVM**, we classify a fruit x as:

$$\hat{y} = \text{sign}(w \cdot x + b)$$

- $w \cdot x$ = **dot product** of the weight vector w and the fruit's features x
- b = bias term (shifts the boundary)

SVM classification for Orange vs Apple



- **Red dots** = Apples (class -1)
- **Orange dots** = Oranges (class +1)
- **Blue solid line** = Decision boundary ($w \cdot x + b = 0$)
- **Blue dashed lines** = Margins ($w \cdot x + b = \pm 1$)
- Points above the solid line are classified as **Orange**, and those below as **Apple**, based on the sign of the dot product $w \cdot x + b$.

The separating line

- The boundary between Apples and Oranges is: $w \cdot x + b = 0$
- Points where: $w \cdot x + b > 0 \rightarrow$ **Orange (+1)**
- Points where: $w \cdot x + b < 0 \rightarrow$ **Apple (-1)**

Farm examples?

- *Identify diseased vs. healthy plants from leaf measurements.*
- *Separate fertile vs. less fertile plots for planting.*
- *Classify soil samples into “needs fertilizer” vs. “good to go.”*
- *Sort harvest quality into “export grade” vs. “local market.”*